

```

In [2]: #importing all the libraries
import warnings;
#disable warnings
warnings.simplefilter('ignore')

import os
os.environ["KERAS_BACKEND"] = "tensorflow"

#set image data format to channels_last
from tensorflow.keras.backend import set_image_data_format
set_image_data_format('channels_last')

import numpy as np
from keras.datasets import mnist
from matplotlib import pyplot as plt
%matplotlib inline

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.preprocessing.image import ImageDataGenerator

#from keras.utils import np_utils
#from tensorflow.keras.utils import np_utils

#np_utils no longer exists in this package
#replace with to_categorical, should be equivalent
from tensorflow.keras.utils import to_categorical

```

WARNING:tensorflow:From C:\Users\vbque\anaconda3\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```

In [3]: #read in the data
data = r'C:\Users\vbque\OneDrive\Desktop\Machine Learning\Assignment6\Assignment 6\data

#data generators
batch_size = 32
img_height, img_width = 28, 28

train_datagen = ImageDataGenerator(rescale=1./255, validation_split=0.2)

#train
train_generator = train_datagen.flow_from_directory(
    data,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical',
    subset='training'
)

```

Found 60 images belonging to 2 classes.

```

In [4]: #validate
validation_generator = train_datagen.flow_from_directory(
    data,
    target_size=(img_height, img_width),

```

```

    batch_size=batch_size,
    class_mode='categorical',
    subset='validation'
)

```

Found 14 images belonging to 2 classes.

```

In [5]: #make CNN model
model = Sequential()

model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 3)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())

#model.add(Dense(128, activation='relu'))
#model.add(Dropout(0.5))
#model.add(Dense(10, activation='softmax'))

model.add(Dense(64, activation='relu'))
model.add(Dense(2, activation='softmax')) #2 classes: covid and normal

model.summary()

```

WARNING:tensorflow:From C:\Users\vbque\anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\vbque\anaconda3\Lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	896
conv2d_1 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 12, 12, 64)	0
dropout (Dropout)	(None, 12, 12, 64)	0
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 64)	589888
dense_1 (Dense)	(None, 2)	130
=====		
Total params: 609410 (2.32 MB)		
Trainable params: 609410 (2.32 MB)		
Non-trainable params: 0 (0.00 Byte)		

```

In [6]: #compile
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

WARNING:tensorflow:From C:\Users\vbque\anaconda3\Lib\site-packages\keras\src\optimize_rs__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

In [7]: *#train the model*

```
epochs = 10

history = model.fit(train_generator,
                    epochs=epochs,
                    validation_data=validation_generator,
                    steps_per_epoch=len(train_generator),
                    validation_steps=len(validation_generator))
```

Epoch 1/10

WARNING:tensorflow:From C:\Users\vbque\anaconda3\Lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\vbque\anaconda3\Lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

2/2 [=====] - 2s 707ms/step - loss: 0.6550 - accuracy: 0.5167 - val_loss: 0.8210 - val_accuracy: 0.5000

Epoch 2/10

2/2 [=====] - 1s 431ms/step - loss: 0.6739 - accuracy: 0.5333 - val_loss: 0.6068 - val_accuracy: 0.5714

Epoch 3/10

2/2 [=====] - 1s 408ms/step - loss: 0.5011 - accuracy: 0.8167 - val_loss: 0.5707 - val_accuracy: 0.5714

Epoch 4/10

2/2 [=====] - 1s 485ms/step - loss: 0.4277 - accuracy: 0.8833 - val_loss: 0.5167 - val_accuracy: 0.7857

Epoch 5/10

2/2 [=====] - 1s 505ms/step - loss: 0.2941 - accuracy: 0.9667 - val_loss: 0.4703 - val_accuracy: 0.7857

Epoch 6/10

2/2 [=====] - 1s 391ms/step - loss: 0.2173 - accuracy: 0.9833 - val_loss: 0.4446 - val_accuracy: 0.7857

Epoch 7/10

2/2 [=====] - 1s 543ms/step - loss: 0.1409 - accuracy: 0.9667 - val_loss: 0.4155 - val_accuracy: 0.8571

Epoch 8/10

2/2 [=====] - 1s 338ms/step - loss: 0.0959 - accuracy: 0.9833 - val_loss: 0.4747 - val_accuracy: 0.8571

Epoch 9/10

2/2 [=====] - 1s 567ms/step - loss: 0.0526 - accuracy: 1.0000 - val_loss: 0.4268 - val_accuracy: 0.9286

Epoch 10/10

2/2 [=====] - 1s 532ms/step - loss: 0.0593 - accuracy: 0.9833 - val_loss: 0.4794 - val_accuracy: 0.9286

In [8]: *#model evaluation*

```
score = model.evaluate(validation_generator, verbose=1)
print(f'Test loss: {score[0]}, Test accuracy: {score[1]}')

#test accuracy is 92.86%
```

1/1 [=====] - 0s 161ms/step - loss: 0.4794 - accuracy: 0.9286

Test loss: 0.4794378876686096, Test accuracy: 0.9285714030265808

```
In [9]: #test set predictions
y_hat = model.predict(validation_generator)
```

1/1 [=====] - 0s 199ms/step

```
In [ ]: #!pip install opencv-python
```

```
In [28]: import matplotlib.pyplot as plt
import numpy as np

plt.figure(figsize=(20, 8)) #increase the figure size for better visualization
j = 1
max_plots = 20 #20 images to be predicted
break_outer_loop = False
class_labels = {0: 'COVID', 1: 'Normal'} #0 is COVID and 1 is Normal

for batch in validation_generator:
    images, labels = batch

    y_hat = model.predict(images)

    for i in range(len(y_hat)):
        threshold1 = y_hat[i] < 0.1
        y_hat[i][threshold1] = 0
        threshold2 = y_hat[i] > 0.9
        y_hat[i][threshold2] = 1
        y_hat[i] = y_hat[i].astype(int)

        estimated = np.argmax(y_hat[i])
        true = np.argmax(labels[i])

        plt.subplot(4, 5, j) #4 rows, 5 images per row, 20 pics total
        title = f'True: {class_labels[true]}, Estimated: {class_labels[estimated]}'
        plt.imshow(images[i], interpolation='bilinear') #makes images clearer
        plt.title(title)

        if estimated == true:
            correct_predictions += 1

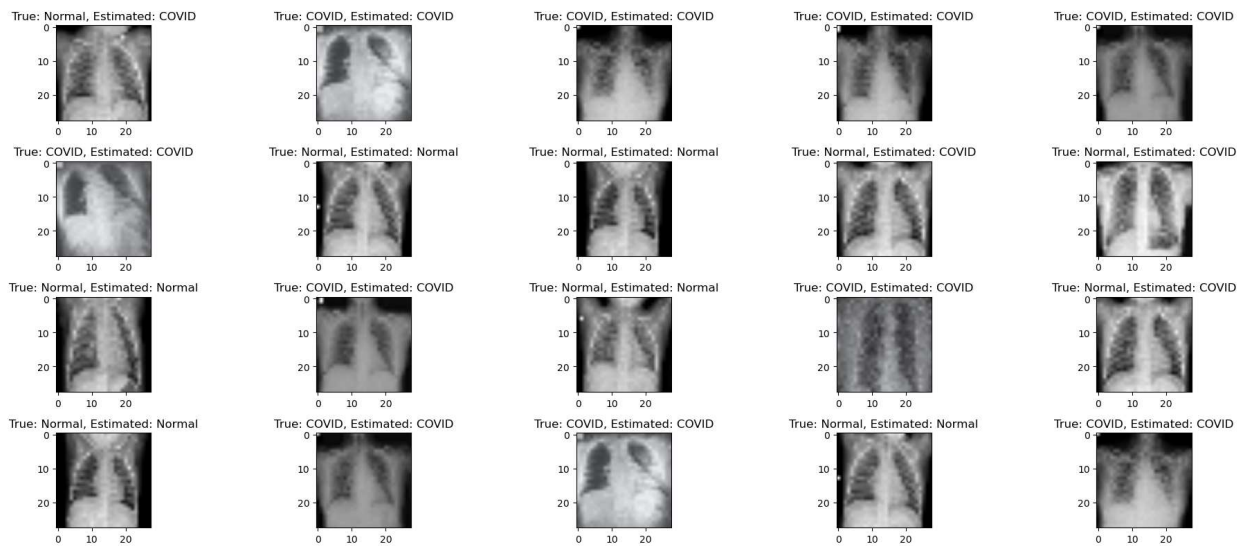
        j += 1
        if j > max_plots:
            break_outer_loop = True
            break

    if break_outer_loop:
        break

plt.tight_layout()
plt.show()
```

1/1 [=====] - 0s 14ms/step

1/1 [=====] - 0s 15ms/step



We can see in the 20 images in the output above (different each run), the classification under True: Normal or True: COVID shows the true classification of the images (which folder they come from). The Estimated: Normal or Estimated: COVID shows what the classification model predicted them to be.

The calculation of accuracy in one of the earlier chunks shows that the test accuracy is 92.86%.

False negatives occur as True: COVID,
Estimated: Normal

False positives occur as True: Normal,
Estimated: COVID

I ran the model a few times to get various outcomes in classification predictions. In some of the runs, there were a couple false negatives or positives (~86% accurate), but

some runs had 100% correct predictions (20/20).

Ultimately, the overall accuracy of 92.86% is good, and could potentially be increased by further normalizing the data, or having a larger dataset to train and test the model on.

In []: